

Gibbs sampler and MCMC

Jingchen (Monika) Hu

Vassar College

MATH 347 Bayesian Statistics

Outline

- 1 Example: Expenditures in the Consumer Expenditure Surveys
- 2 Prior and posterior distributions for mean AND standard deviation
- 3 Use JAGS (Just Another Gibbs Sampler) and Bayesian inferences
- 4 MCMC diagnostics
- 5 Recap

Outline

- 1 Example: Expenditures in the Consumer Expenditure Surveys
- 2 Prior and posterior distributions for mean AND standard deviation
- 3 Use JAGS (Just Another Gibbs Sampler) and Bayesian inferences
- 4 MCMC diagnostics
- 5 Recap

The TOTEXPPQ variable in the CE sample

```
CEsample <- read_csv("CEsample1.csv")
```

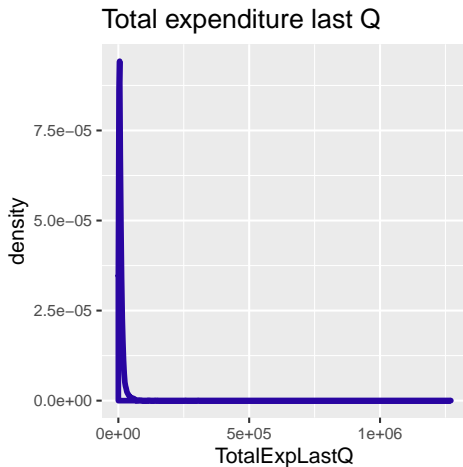
```
summary(CEsample$TotalExpLastQ)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      30    3522    6417    9513   11450  1270598
```

```
sd(CEsample$TotalExpLastQ)
```


```
## [1] 19341.25
```

The TOTEXPPQ variable in the CE sample cont'd



Log transformation of the TOTEXPPQ variable

```
CEsample$LogTotalExpLastQ <- log(CEsample$TotalExpLastQ)
ggplot(data = CEsample, aes(LogTotalExpLastQ)) +
  geom_density(color = crcblue, size = 1) +
  labs(title = "Total expenditure last Q (log)") +
  theme_grey(base_size = 8, base_family = "")
```




The Normal distribution

- The Normal distribution is a symmetric, bell-shaped distribution.
- It has two parameters: mean μ and standard deviation σ .
- The probability density function (pdf) of $\text{Normal}(\mu, \sigma)$ is:

$$f(y) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right), -\infty < y < \infty.$$

The Normal distribution cont'd



i.i.d. Normals

- Suppose there are a sequence of n responses: Y_1, Y_2, \dots, Y_n .
- Further suppose each response **independently and identically** follows a Normal distribution:

$$Y_i \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma).$$

- Then the joint probability density function (joint pdf) of y_1, \dots, y_n is:

$$f(y_1, \dots, y_n) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(\frac{-(y_i - \mu)^2}{2\sigma^2}\right), -\infty < y_i < \infty. \quad (1)$$

Outline

- 1 Example: Expenditures in the Consumer Expenditure Surveys
- 2 Prior and posterior distributions for mean AND standard deviation**
- 3 Use JAGS (Just Another Gibbs Sampler) and Bayesian inferences
- 4 MCMC diagnostics
- 5 Recap

Step 1: Prior distributions

- The data model/sampling density for n observations:

$$Y_i \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma).$$

- There are two parameters μ and σ in the Normal model.
- Therefore, the likelihood is in terms of both unknown parameters:

$$f(y_1, \dots, y_n) = L(\mu, \sigma). \quad (2)$$

Step 1: Prior distributions

- The data model/sampling density for n observations:

$$Y_i \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma).$$

- There are two parameters μ and σ in the Normal model.
- Therefore, the likelihood is in terms of both unknown parameters:

$$f(y_1, \dots, y_n) = L(\mu, \sigma). \quad (2)$$

- Need a joint prior distribution:

$$\pi(\mu, \sigma). \quad (3)$$

- Bayes' rule will help us derive a joint posterior:

$$\pi(\mu, \sigma \mid y_1, \dots, y_n). \quad (4)$$

If only mean μ is unknown: Normal conjugate prior

- For this special case, Normal prior for μ is a conjugate prior:
 - ▶ The prior distribution:

$$\mu \mid \sigma \sim \text{Normal}(\mu_0, \sigma_0). \quad (5)$$

- ▶ The sampling density:

$$y_1, \dots, y_n \mid \mu, \sigma \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma). \quad (6)$$

If only mean μ is unknown: Normal conjugate prior

- For this special case, Normal prior for μ is a conjugate prior:
 - The prior distribution:

$$\mu \mid \sigma \sim \text{Normal}(\mu_0, \sigma_0). \quad (5)$$

- The sampling density:

$$y_1, \dots, y_n \mid \mu, \sigma \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma). \quad (6)$$

- The posterior distribution:

$$\mu \mid y_1, \dots, y_n, \phi \sim \text{Normal} \left(\frac{\phi_0 \mu_0 + n \phi \bar{y}}{\phi_0 + n \phi}, \sqrt{\frac{1}{\phi_0 + n \phi}} \right), \quad (7)$$

where $\phi = \frac{1}{\sigma^2}$ (and $\phi_0 = \frac{1}{\sigma_0^2}$), the precision. Since σ (and σ_0) is known, ϕ (and ϕ_0) is known too.

- We can then use the `rnorm()` R function to sample posterior draws of μ from Equation (7). Known quantities: $\phi_0, \mu_0, n, \bar{y}, \phi$

If only standard deviation σ is unknown: Gamma conjugate prior for $1/\sigma^2$

- For this special case, Gamma prior for $1/\sigma^2$ is a conjugate prior:
 - ▶ The prior distribution:

$$1/\sigma^2 \mid \mu \sim \text{Gamma}(\alpha, \beta). \quad (8)$$

- ▶ The sampling density:

$$y_1, \dots, y_n \mid \mu, \sigma \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma). \quad (9)$$

If only standard deviation σ is unknown: Gamma conjugate prior for $1/\sigma^2$

- For this special case, Gamma prior for $1/\sigma^2$ is a conjugate prior:
 - The prior distribution:

$$1/\sigma^2 \mid \mu \sim \text{Gamma}(\alpha, \beta). \quad (8)$$

- The sampling density:

$$y_1, \dots, y_n \mid \mu, \sigma \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma). \quad (9)$$

- The posterior distribution:

$$1/\sigma^2 \mid y_1, \dots, y_n, \mu \sim \text{Gamma} \left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (y_i - \mu)^2 \right) \quad (10)$$

- We can then use `rgamma()` R function to sample posterior draws of σ from Equation (10). **Known quantities:** $\alpha, n, \beta, \{y_i\}, \mu$

Go back to: both μ and σ are unknown

- Given what we have seen, how about a joint prior distribution as:

$$\pi(\mu, \sigma) = \pi_1(\mu)\pi_2(\sigma), \quad (11)$$

and let priors be

$$\mu \sim \text{Normal}(\mu_0, \sigma_0), \quad (12)$$

$$1/\sigma^2 \sim \text{Gamma}(\alpha, \beta). \quad (13)$$

Go back to: both μ and σ are unknown

- Given what we have seen, how about a joint prior distribution as:

$$\pi(\mu, \sigma) = \pi_1(\mu)\pi_2(\sigma), \quad (11)$$

and let priors be

$$\mu \sim \text{Normal}(\mu_0, \sigma_0), \quad (12)$$

$$1/\sigma^2 \sim \text{Gamma}(\alpha, \beta). \quad (13)$$

Bayes' rule will produce two **full conditional posterior distributions**:

$$\mu \mid y_1, \dots, y_n, \phi \sim \text{Normal} \left(\frac{\phi_0 \mu_0 + n \phi \bar{y}}{\phi_0 + n \phi}, \sqrt{\frac{1}{\phi_0 + n \phi}} \right), \quad (14)$$

$$1/\sigma^2 = \phi \mid y_1, \dots, y_n, \mu \sim \text{Gamma} \left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (y_i - \mu)^2 \right). \quad (15)$$

Go back to: both μ and σ are unknown

- Given what we have seen, how about a joint prior distribution as:

$$\pi(\mu, \sigma) = \pi_1(\mu)\pi_2(\sigma), \text{ requires independence} \quad (11)$$

and let priors be

$$\mu \sim \text{Normal}(\mu_0, \sigma_0), \quad (12)$$

$$1/\sigma^2 \sim \text{Gamma}(\alpha, \beta). \quad (13)$$

Bayes' rule will produce two **full conditional posterior distributions**:

$$\mu \mid y_1, \dots, y_n, \phi \sim \text{Normal} \left(\frac{\phi_0 \mu_0 + n \phi \bar{y}}{\phi_0 + n \phi}, \sqrt{\frac{1}{\phi_0 + n \phi}} \right), \quad (14)$$

even though unknown, we assume known

$$1/\sigma^2 = \phi \mid y_1, \dots, y_n, \mu \sim \text{Gamma} \left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (y_i - \mu)^2 \right). \quad (15)$$

- Why and how to derive them? Key: independent priors.

Sampling scheme: A Gibbs sampler

$$\begin{aligned}\mu \mid y_1, \dots, y_n, \phi &\sim \text{Normal} \left(\frac{\phi_0 \mu_0 + n \phi \bar{y}}{\phi_0 + n \phi}, \sqrt{\frac{1}{\phi_0 + n \phi}} \right), \\ 1/\sigma^2 = \phi \mid y_1, \dots, y_n, \mu &\sim \text{Gamma} \left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (y_i - \mu)^2 \right).\end{aligned}$$

Sampling scheme: A Gibbs sampler

$$\mu \mid y_1, \dots, y_n, \phi \sim \text{Normal} \left(\frac{\phi_0 \mu_0 + n \phi \bar{y}}{\phi_0 + n \phi}, \sqrt{\frac{1}{\phi_0 + n \phi}} \right),$$

$$1/\sigma^2 = \phi \mid y_1, \dots, y_n, \mu \sim \text{Gamma} \left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (y_i - \mu)^2 \right).$$

Start with initial values for parameters, $(\mu^{(0)}, \phi^{(0)})$. For $s = 1, \dots, S$, generate from the following sequence of **full conditional posterior distributions**:

$$- \mu^{(s)} \sim p(\mu \mid \phi^{(s-1)}, Y)$$

essentially Markov Chain

$$- \phi^{(s)} \sim p(\phi \mid \mu^{(s)}, Y)$$

you can find that μ_0 is never used!

$$- \text{Set } \theta^{(s)} = (\mu^{(s)}, \phi^{(s)})$$

The sequence $\{\theta^{(s)}: s = 1, \dots, S\}$ may be viewed (but is not necessarily...yet) as a dependent sample from the joint posterior distribution of (μ, ϕ) .

Use R/RStudio to run a Gibbs sampler

```
gibbs_normal <- function(input, S, seed){
  set.seed(seed)
  ybar <- mean(input$y)
  n <- length(input$y)
  para <- matrix(0, S, 2)  # 2xS
  phi <- input$phi_init
  for(s in 1:S){
    mu1 <- (input$mu_0/input$sigma_0^2 + n*phi*ybar)/
      (1/input$sigma_0^2 + n*phi)
    sigma1 <- sqrt(1/(1/input$sigma_0^2 + n*phi))
    mu <- rnorm(1, mean = mu1, sd = sigma1)
    alpha1 <- input$alpha + n/2
    beta1 <- input$beta + sum((input$y - mu)^2)/2
    phi <- rgamma(1, shape = alpha1, rate = beta1)
    para[s, ] <- c(mu, phi)
  }
  para }
```

Use R/RStudio to run a Gibbs sampler cont'd

- Run the Gibbs sampler:

```
input <- list(y = CEsample$LogTotalExpLastQ, mu_0 = 5, sigma_0 = 1,
alpha = 1, beta = 1, phi_init = 1)
output <- gibbs_normal(input, S = 10000, seed = 123)
```

Use R/RStudio to run a Gibbs sampler cont'd

- Run the Gibbs sampler:


```
input <- list(y = CEsample$LogTotalExpLastQ, mu_0 = 5, sigma_0 = 1,
alpha = 1, beta = 1, phi_init = 1)
output <- gibbs_normal(input, S = 10000, seed = 123)
```

- Extract posterior draws of μ and ϕ from the Gibbs sampler output:

```
para_post <- as.data.frame(output)
names(para_post) <- c("mu", "phi")
```


Use R/RStudio to run a Gibbs sampler cont'd

```
ggplot(para_post, aes(mu)) +
  geom_density(size = 2, color = "blue") +
  labs(title = "Posterior draws of mu") +
  theme_grey(base_size = 8,
    base_family = "")
```




```
quantile(para_post$mu, c(0.025, 0.975))
```

```
##      2.5%      97.5%
## 8.725482 8.769543
```

Use R/RStudio to run a Gibbs sampler cont'd

```
ggplot(para_post, aes(phi)) +
  geom_density(size = 2, color = "blue") +
  labs(title = "Posterior draws of phi") +
  theme_grey(base_size = 8,
    base_family = "")
```



```
quantile(para_post$phi, c(0.025,0.975))
```

```
##      2.5%      97.5%
## 1.206482 1.294191
```

Exercise: Gibbs sampler

- Exercise 1: Update the following Gibbs sampler to initiate it with `mu_init`. Compare your results to the existing results.
- Exercise 2: Update the following Gibbs sampler to initiate it with `mu_init` and `phi_init`. Compare your results to the existing results.

Extends to more than 2 parameters

Suppose the full conditional posterior distributions of $\theta = (\theta_1, \theta_2, \dots, \theta_m)$ are all tractable.

- Start the Gibbs sampler at the initial value $(\theta_1^{(0)}, \theta_2^{(0)}, \dots, \theta_m^{(0)})$, or a subset of them.
- For the $(s + 1)$ -th iteration: generate from the following full conditionals sequentially:

with m parameters, we need $m-1$ initial values for each of the other parameters to get the Gibbs sampling works

- ▶ $\theta_1^{(s+1)} \sim p(\theta_1 \mid \theta_2^{(s)}, \theta_3^{(s)}, \theta_4^{(s)}, \dots, \theta_m^{(s)}, Y)$
- ▶ $\theta_2^{(s+1)} \sim p(\theta_2 \mid \theta_1^{(s+1)}, \theta_3^{(s)}, \theta_4^{(s)}, \dots, \theta_m^{(s)}, Y)$
- ▶ $\theta_3^{(s+1)} \sim p(\theta_3 \mid \theta_1^{(s+1)}, \theta_2^{(s+1)}, \theta_4^{(s)}, \dots, \theta_m^{(s)}, Y)$
- ▶ \vdots
- ▶ $\theta_m^{(s+1)} \sim p(\theta_m \mid \theta_1^{(s+1)}, \theta_2^{(s+1)}, \theta_3^{(s+1)}, \dots, \theta_{m-1}^{(s+1)}, Y)$

Outline

- 1 Example: Expenditures in the Consumer Expenditure Surveys
- 2 Prior and posterior distributions for mean AND standard deviation
- 3 Use JAGS (Just Another Gibbs Sampler) and Bayesian inferences**
- 4 MCMC diagnostics
- 5 Recap

JAGS for unknown mean and standard deviation case

- R package `runjags` to run Markov chain Monte Carlo simulations.
- Descriptive of the sampling model and the prior.

JAGS for unknown mean and standard deviation case

- R package `runjags` to run Markov chain Monte Carlo simulations.
- Descriptive of the sampling model and the prior.
- Installing JAGS software and `runjags` R package
 - ▶ Download JAGS at [this link](#)
 - ▶ Install and load `runjags` R package

```
install.packages("runjags")
```

```
library(runjags)
```

JAGS for unknown mean and standard deviation case cont'd

- Example: Normal model with unknown mean and standard deviation.
 - ▶ The prior distributions:

$$\begin{aligned}\mu &\sim \text{Normal}(\mu_0, \sigma_0), \\ 1/\sigma^2 = \phi &\sim \text{Gamma}(\alpha, \beta).\end{aligned}$$

- ▶ The sampling density:

$$y_1, \dots, y_n \mid \mu, \sigma \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma).$$

- ▶ Bayes' rule will produce two **full conditional posterior distributions**:

$$\begin{aligned}\mu \mid y_1, \dots, y_n, \phi &\sim \text{Normal}\left(\frac{\phi_0 \mu_0 + n \bar{y} \phi}{\phi_0 + n \phi}, \sqrt{\frac{1}{\phi_0 + n \phi}}\right), \\ 1/\sigma^2 = \phi \mid y_1, \dots, y_n, \mu &\sim \text{Gamma}\left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (y_i - \mu)^2\right).\end{aligned}$$

JAGS for unknown mean and standard deviation case cont'd

- Only need to focus on the sampling density and the prior:
 - The sampling density:

$$y_1, \dots, y_n \mid \mu, \sigma \stackrel{i.i.d.}{\sim} \text{Normal}(\mu, \sigma).$$

- The prior distributions:

$$\begin{aligned} \mu &\sim \text{Normal}(\mu_0, \sigma_0), \\ 1/\sigma^2 = \phi &\sim \text{Gamma}(\alpha, \beta). \end{aligned}$$

```
modelString <- "
model{
  for (i in 1:N) {
    y[i] ~ dnorm(mu, phi)
  }
  mu ~ dnorm(mu_0, phi_0)
  phi ~ dgamma(alpha, beta)
}
```

JAGS for unknown mean and standard deviation case cont'd

- Pass the data and hyperparameter values to JAGS:

```
y <- CEsample$LogTotalExpLastQ
N <- length(y)
the_data <- list("y" = y, "N" = N, "mu_0" = 5, "phi_0" = 1/1^2,
"alpha" = 1, "beta" = 1)
```

JAGS for unknown mean and standard deviation case cont'd

hyper-param=param of param

- Pass the data and hyperparameter values to JAGS:

```
y <- CEsample$LogTotalExpLastQ
N <- length(y)
the_data <- list("y" = y, "N" = N, "mu_0" = 5, "phi_0" = 1/1^2,
"alpha" = 1,"beta" = 1)
```

$\phi_0 = 1/\sigma_0^2$

- Run the JAGS code for this model:

```
posterior <- run.jags(modelString,
  data = the_data,
  monitor = c("mu", "phi"),
  n.chains = 1,
  adapt = 1000,
  burnin = 2000,
  sample = 5000,
  thin = 1)
```

JAGS for unknown mean and standard deviation case

cont'd

- Obtain posterior summaries of μ and ϕ :

```
summary(posterior)
```

##	Lower95	Median	Upper95	Mean	SD	Mode	MCerr	MC%ofSD
## mu	8.72529	8.747810	8.76868	8.747781	0.01111673	NA	0.0001572144	1.4
## phi	1.20652	1.251015	1.29268	1.250759	0.02206280	NA	0.0003091955	1.4
##	SSEff	AC.10	psrf					
## mu	5000	0.030116334	NA					
## phi	5092	-0.002062701	NA					

Outline

- 1 Example: Expenditures in the Consumer Expenditure Surveys
- 2 Prior and posterior distributions for mean AND standard deviation
- 3 Use JAGS (Just Another Gibbs Sampler) and Bayesian inferences
- 4 MCMC diagnostics**
- 5 Recap

MCMC Convergence to posterior distribution

Theory proves that if a Gibbs sampler iterates enough, the draws will be from the joint posterior distribution (called the **target** or **stationary** distribution).

MCMC Convergence to posterior distribution

Theory proves that if a Gibbs sampler iterates enough, the draws will be from the joint posterior distribution (called the **target** or **stationary** distribution).

- Do initial values matter? Should they matter?
- “Markov chain” indicate dependence of draws. How to create independent parameter draws?
- How long do we need to run the MCMC to adequately explore the posterior distribution?
- How can we tell if the chain is not converging?

Discard initial steps: burn-in

- Do not want pre-convergence values to influence the summary of the posterior distribution (the draws of the parameters from the MCMC) too much.
- Ideal is to throw away all draws before convergence to target distribution, and use only draws after convergence.
- Hard to know exactly when convergence has happened. Thus, it is standard to throw out the first $p\%$ of draws (default is 50%) as **burn-in** after you are satisfied that convergence has been achieved.
- Inference done with remaining $(100 - p\%)$ of draws.

Discard initial steps: burn-in cont'd

```
posterior <- run.jags(modelString,  
  data = the_data,  
  monitor = c("mu", "phi"),  
  n.chains = 1,  
  adapt = 1000,  
  burnin = 2000,  
  sample = 5000,  
  thin = 1)
```


Trace plots to check MCMC mixing

- Trace plots show the progress of the chain over time. The ideal trace plot has a lot of movement around the mode from iteration to iteration and not a lot of “stickiness.”
- Trace plots with a clear trend are evidence that MCMC sampler has not yet converged, and it needs to run longer (or there is an error in your code).
- Trace plots with extreme stickiness suggest that the chain may not have explored the entire parameter space yet, and it needs to run longer.
- Hard to tell convergence from trace plots alone, because you might be stuck in one region of the parameter space in all of the iterations.




Trace plots example

```
plot(posterior, vars = "mu")
```

```
## Generating plots...
```



adapt+burn-in=3000







ACF plots and effective sample size to check correlation diagnostics

- Independent (Monte Carlo) draws have zero autocorrelations. Dependent draws (MCMC) have non-zero, usually positive, autocorrelations.
- Large autocorrelations indicate that chain is not mixing well, i.e., the chain possibly has not explored the full space of the posterior distribution (has not converged).
- Large autocorrelations often arise in multi-parameter MCMC algorithms when the parameters are highly correlated.
- This usually means you need to run the chains longer to feel comfortable about convergence. It also implies a smaller effective sample size. . . .
- ACF plots (auto correlation function) & effective sample size.

ACF plots example

```
plot(posterior, vars = "mu")
```

```
## Generating plots...
```



Effective sample size example

how many equivalent monte carlo draws (which are independent) that my MCMC draws represent

- The column of SSeff; recall sample is 5000.

```
summary(posterior)
```

	Lower95	Median	Upper95	Mean	SD	Mode	MCerr	MC%ofSD
## mu	8.72529	8.747810	8.76868	8.747781	0.01111673	NA	0.0001572144	1.4
## phi	1.20652	1.251015	1.29268	1.250759	0.02206280	NA	0.0003091955	1.4
##	SSeff	AC.10	psrf					
## mu	5000	0.030116334	NA					
## phi	5092	-0.002062701	NA					

sometimes, b/c of negative auto-correlation, it leads to effective sample size > 5000, we don't like this either

Create independent samples: thin

- Monte Carlo samples come from independent parameter draws, thus they move around the parameter space freely.
- Markov chain properties could create dependence among adjacent parameter draws.
- Use **thin** to save every x draws.





say you see the ACF plot is still very high
every 50 draws, then you'd like to set
thin=50

```
posterior <- run.jags(modelString,
  data = the_data,
  monitor = c("mu", "phi"),
  n.chains = 1,
  adapt = 1000,
  burnin = 2000,
  sample = 5000,
  thin = 1)
```

MCMC diagnostics for the CE example

```
plot(posterior, vars = "mu")
```





```
## Generating plots...
```



MCMC diagnostics for the CE example cont'd

```
plot(posterior, vars = "phi")
```

```
## Generating plots...
```



Gelman-Rubin diagnostics

- Gelman and Rubin (1992) propose running $m > 1$ separate Markov chains each of length S . When convergence is reached, the draws from the chains are all from the target distribution. Hence, at convergence, the output from all chains is indistinguishable.
- Diagnostic based on the contrapositive: if output from chains is not indistinguishable, convergence has not been reached. Diagnostic applied to one variable at a time.
- Use different starting values that are overdispersed (really spread out) relative to the posterior distribution. This helps ensure all the chains are not stuck in the same region.
- Gelman-Rubin diagnostic based a comparison of within-chain and between-chain variances, which is similar to a classical analysis of variance.

Gelman-Rubin diagnostics example

- Create intinial values of mu and phi:

```
inits1 <- dump.format(list(mu = 1, phi = 1,  
  .RNG.name="base::Super-Duper", .RNG.seed = 1))  
inits2 <- dump.format(list(mu = 10, phi = 10,  
  .RNG.name="base::Wichmann-Hill", .RNG.seed = 2))
```

Gelman-Rubin diagnostics example

- Create initial values of μ and ϕ :

```
inits1 <- dump.format(list(mu = 1, phi = 1,
  .RNG.name="base::Super-Duper", .RNG.seed = 1))
inits2 <- dump.format(list(mu = 10, phi = 10,
  .RNG.name="base::Wichmann-Hill", .RNG.seed = 2))
```

- Feed in `inits1` and `inits2`, and let `n.chains = 2`:

```
posterior_2chains <- run.jags(modelString,
  data = the_data,
  monitor = c("mu", "phi"),
  n.chains = 2,
  inits=c(inits1, inits2),
  adapt = 1000,
  burnin = 2000,
  sample = 5000,
  thin = 1)
```

running two chains

Gelman-Rubin diagnostics example cont'd

- Return `psrf` from the output, as Gelman-Rubin diagnostic results:





```
posterior_2chains$psrf
```

```
## Potential scale reduction factors:
##
##      Point est. Upper C.I.
## mu           1           1
## phi          1           1
##
## Multivariate psrf (for all monitored variables):
##
## 1
##
## Target psrf
##
## 1.05
```

MCMC diagnostics for the CE example. 2 chains

```
plot(posterior_2chains, vars = "mu")
```

```
## Generating plots...
```



Useful diagnostics/functions in coda package

The coda R package provides many popular diagnostics for assessing convergence of MCMC output from self-coded Gibbs sampler.

- Trace plots: `traceplot()`
- Autocorrelation: `autocorr.plot()`
- Effective Sample Size: `effectiveSize()`
- Gelman-Rubin: `gelman.diag()`

Useful diagnostics/functions in coda package

- One needs to convert parameter draws into an MCMC object. For example:

```
install.packages("coda")
```

```
library(coda)
```

```
out <- gibbs_normal(input, S = 10000, seed = 123)
```

```
para_post = as.data.frame(out)
```

```
names(para_post) = c("mu", "phi")
```


Useful diagnostics/functions in coda package

- One needs to convert parameter draws into an MCMC object. For example:

```
install.packages("coda")
```

```
library(coda)
```

```
out <- gibbs_normal(input, S = 10000, seed = 123)
```

```
para_post = as.data.frame(out)
```

```
names(para_post) = c("mu", "phi")
```

- Then one can perform MCMC diagnostics. For example:

```
mu.mcmc = as.mcmc(para_post$mu)
```

```
traceplot(mu.mcmc)
```

```
autocorr.plot(mu.mcmc)
```

```
effectiveSize(mu.mcmc)
```

```
gelman.diag(mu.mcmc)
```

Note: `gelman.diag()` needs at least 2 chains.

Outline

- 1 Example: Expenditures in the Consumer Expenditure Surveys
- 2 Prior and posterior distributions for mean AND standard deviation
- 3 Use JAGS (Just Another Gibbs Sampler) and Bayesian inferences
- 4 MCMC diagnostics
- 5 Recap**

Recap

- Bayesian inference procedure:
 - ▶ Step 1: express an opinion about the location of mean μ and standard deviation σ (or precision ϕ) before sampling (prior).
 - ▶ Step 2: take the sample (data/likelihood).
 - ▶ Step 3: use Bayes' rule to sharpen and update the previous opinion about μ and σ (or precision ϕ) given the information from the sample (posterior).

Recap

- Bayesian inference procedure:
 - ▶ Step 1: express an opinion about the location of mean μ and standard deviation σ (or precision ϕ) before sampling (prior).
 - ▶ Step 2: take the sample (data/likelihood).
 - ▶ Step 3: use Bayes' rule to sharpen and update the previous opinion about μ and σ (or precision ϕ) given the information from the sample (posterior).
- For Normal data/likelihood, Normal distributions are conjugate priors for μ , and Gamma distributions are conjugate priors for ϕ . The same goes for the full conditional posterior distributions.

Recap

- Bayesian inference procedure:
 - ▶ Step 1: express an opinion about the location of mean μ and standard deviation σ (or precision ϕ) before sampling (prior).
 - ▶ Step 2: take the sample (data/likelihood).
 - ▶ Step 3: use Bayes' rule to sharpen and update the previous opinion about μ and σ (or precision ϕ) given the information from the sample (posterior).
- For Normal data/likelihood, Normal distributions are conjugate priors for μ , and Gamma distributions are conjugate priors for ϕ . The same goes for the full conditional posterior distributions.
- One can either write a loop function to run a Gibbs sampler, or use JAGS.

Recap cont'd

- MCMC diagnostics!
 - ▶ Diagnostics cannot guarantee that chain has converged.
 - ▶ Diagnostics can indicate that it has not converged.
 - ▶ Solutions: run longer, thinning, standardize variables in multi-parameter models (can reduce auto correlations).